

modeling_reg

March 16, 2023

```
[1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import statsmodels.api as sm
import warnings
import xgboost as xgb
from statsmodels.stats.anova import anova_lm
from sklearn.decomposition import PCA
from sklearn.preprocessing import StandardScaler
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn.model_selection import KFold
from sklearn.linear_model import LinearRegression
from sklearn.model_selection import GridSearchCV
from sklearn.model_selection import cross_val_score
from statsmodels.gam.api import GLMGam, BSplines

# import diagnostics from smlearn website
import SM_plots as SM

warnings.filterwarnings('ignore')
```

```
[2]: # load data
path = '/Users/dylanjorling/UCLA/412proj/data/'
name = 'full_combine_data'
data = pd.read_csv(path + name, index_col=0)
data.shape
```

[2]: (13210, 15)

```
[3]: ##### EDA Plots #####
combine_complete = data.dropna()
print(combine_complete.shape)
combine_complete.head()
combine_complete = combine_complete[combine_complete['pick'] != 'undrafted']
```

```

#combine_complete = combine_complete[combine_complete['pos'] == 'LB'] #filter
↳for specific positions
#combine_complete = combine_complete[combine_complete['year'] >= 2010] #filter
↳for specific year...didn't alter much
combine_complete['pick'] = combine_complete['pick'].astype(int)
combine_complete = combine_complete[combine_complete['pick'] < 300]
combine_complete = combine_complete[combine_complete['bench'] < 74]
pos_one_hot = combine_complete['pos'].str.get_dummies()

y = combine_complete['pick']
X = combine_complete.iloc[:, 4:-1]
X = pd.concat([X, pos_one_hot], axis=1)
null_sd = y.std()
print(null_sd)

combine_complete.head()

```

```

(5843, 15)
69.80294730041324

```

```

[3]:
   year  name      college pos  height  weight  hand_size \
3350  1997  John Allred      USC  TE    76.4    244    10.00
3355  1997  Duane Ashman    Virginia  DE    75.5    274     9.50
3357  1997  Raymond Austin  Tennessee  DB    70.6    190    10.25
3360  1997  Antonio Banks  Virginia Tech  DB    70.0    203     9.50
3361  1997  Ronde Barber    Virginia  DB    69.4    185     9.50

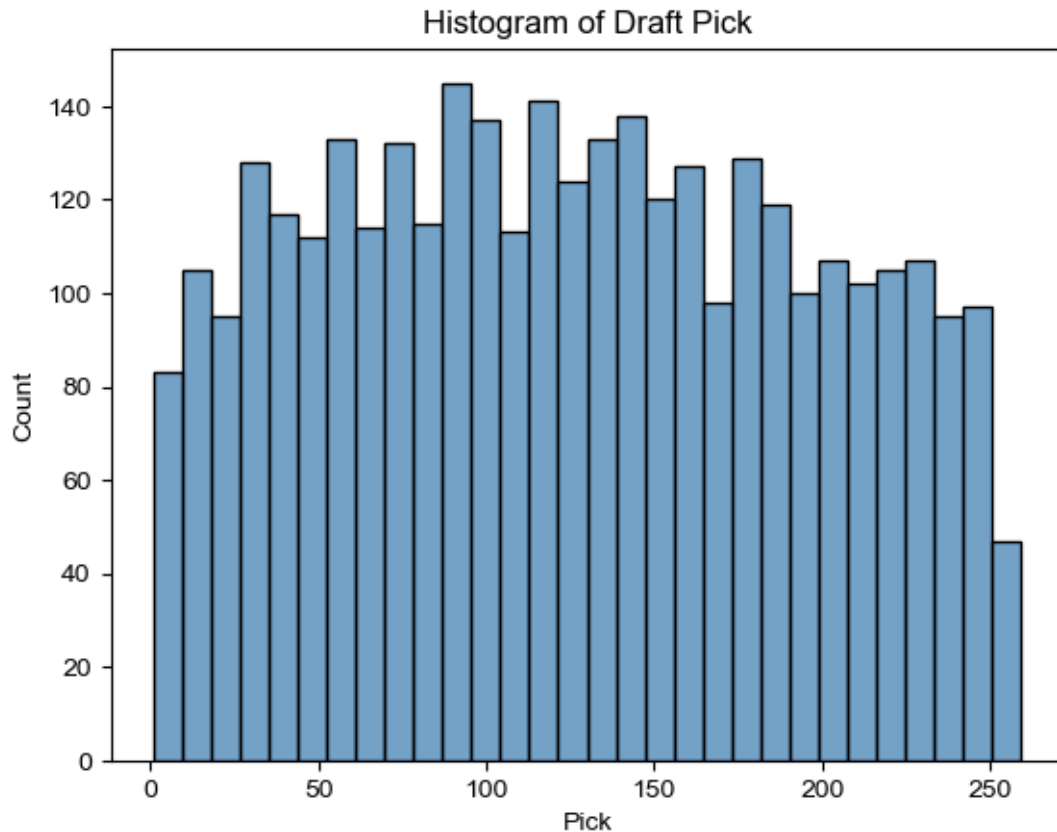
   arm_length  forty  bench  vert  broad_jump  shuttle  3cone  pick
3350      32.50   5.01   15.0  32.0      112.0     4.32   7.45   38
3355      35.63   5.03   19.0  32.5      110.0     4.83   8.80  161
3357      32.00   4.80   12.0  32.5      111.0     4.05   7.14  145
3360      32.00   4.66   18.0  36.0      117.0     4.41   7.85  113
3361      31.63   4.68   14.0  34.5      118.0     4.46   7.22   66

```

```

[5]: ax = sns.histplot(combine_complete, x='pick', color="steelblue", bins=30)
ax.set(xlabel='Pick', ylabel='Count', title='Histogram of Draft Pick')
sns.set_style("whitegrid")
plt.show()
# underlying distribution is uniform

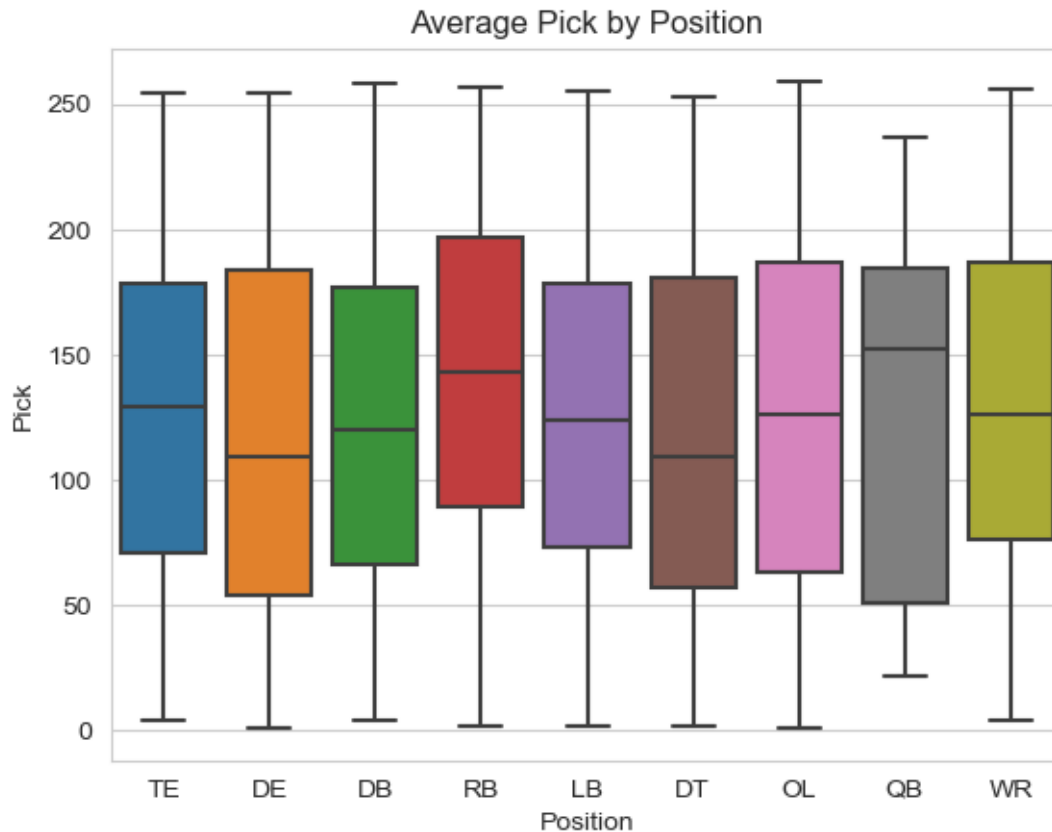
```



```
[6]: pos_means = combine_complete.groupby('pos').mean()['pick']
ax = sns.boxplot(data=combine_completex=pos_means, y=pos_means.index, orient='h')
ax.set(xlabel='Pick', ylabel='Position', title='Average Pick by Position')
sns.set_style("whitegrid")
plt.show()
```

```
File "/var/folders/z7/1q6khp0n3xq0p_97cq1hvpd40000gn/T/ipykernel_2929/2923876733.
↳py", line 2
    ax = sns.boxplot(data=combine_completex=pos_means, y=pos_means.index,
↳orient='h')
                        ^
SyntaxError: invalid syntax
```

```
[7]: pos_means = combine_complete.groupby('pos').mean()['pick']
ax = sns.boxplot(data=combine_complete, x='pos', y='pick')
ax.set(xlabel='Position', ylabel='Pick', title='Average Pick by Position')
sns.set_style("whitegrid")
plt.show()
```



[101]: *# This is a plot of some of the more predictive variables*

```
fig, axs = plt.subplots(nrows=2, ncols=2)
sns.scatterplot(data=combine_complete,
                x='bench',
                y='weight',
                hue='pos',
                s=3,
                ax=axs[0,0],
                legend=False)
axs[0, 0].set(xlabel='Bench Press Reps', ylabel='Weight', title='Bench vs_
↳Weight')
sns.regplot(data=combine_complete,
            x='bench',
            y='weight',
            scatter=False,
            ci=None,
            color="navy",
            ax=axs[0,0])
sns.set_style("whitegrid")
```

```

sns.scatterplot(data=combine_complete,
                x='vert',
                y='broad_jump',
                hue='pos',
                s=3,
                ax=axes[0,1],
                legend=False)
axes[0, 1].set(xlabel='Vert', ylabel='Broad Jump', title='Vertical vs. Broad_
↳ Jump')
sns.regplot(data=combine_complete,
            x='vert',
            y='broad_jump',
            scatter=False,
            ci=None,
            color="navy",
            ax=axes[0,1])
sns.set_style("whitegrid")

sns.scatterplot(data=combine_complete,
                x='forty',
                y='shuttle',
                hue='pos',
                s=3,
                ax=axes[1,0],
                legend=False)
axes[1, 0].set(xlabel='Forty', ylabel='Shuttle', title='Forty vs. Shuttle')
sns.regplot(data=combine_complete,
            x='forty',
            y='shuttle',
            scatter=False,
            ci=None,
            color="navy",
            ax=axes[1,0])
sns.set_style("whitegrid")

sns.scatterplot(data=combine_complete,
                x='hand_size',
                y='arm_length',
                hue='pos',
                s=3,
                ax=axes[1,1],
                legend=False)
axes[1, 1].set(xlabel='Hand Size', ylabel='Arm Length', title='Hand Size vs. Arm_
↳ Length')

sns.regplot(data=combine_complete,

```

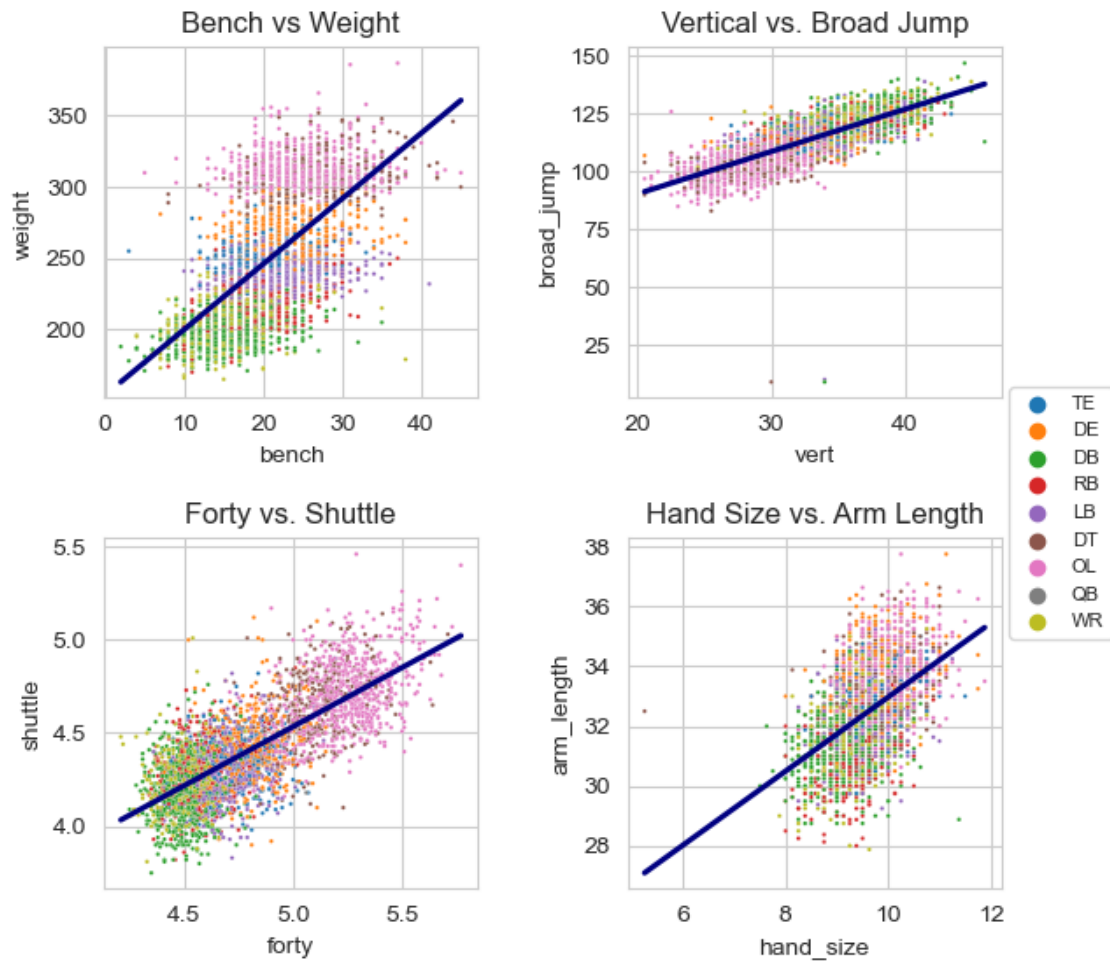
```

x='hand_size',
y='arm_length',
scatter=False,
ci=None,
color="navy",
ax=axes[1,1])
sns.set_style("whitegrid")

#handles, labels = axes[1,1].get_legend_handles_labels()
fig.legend(handles=handles, labels=labels, loc='center right', fontsize = 8)
plt.gca().get_legend_handles_labels()
plt.subplots_adjust(left=0.1,
                    bottom=0.5,
                    right=0.9,
                    top=1.5,
                    wspace=0.4,
                    hspace=0.4)

fig.show()

```



```

[94]: # This is a plot of some of the more predictive variables

fig, axs = plt.subplots(nrows=2, ncols=2)
sns.scatterplot(data=combine_complete,
                x='forty',
                y='pick',
                hue='pos',
                s=3,
                ax=axs[0,0],
                legend=False)
axs[0, 0].set(xlabel='forty yard dash time', ylabel='Pick', title='Forty vs Pick')
sns.regplot(data=combine_complete,
            x='forty',
            y='pick',
            scatter=False,
            ci=None,
            color="navy",
            ax=axs[0,0])
sns.set_style("whitegrid")

sns.scatterplot(data=combine_complete,
                x='weight',
                y='pick',
                hue='pos',
                s=3,
                ax=axs[0,1],
                legend=False)
axs[0, 1].set(xlabel='weight', ylabel='Pick', title='Weight vs Pick')
sns.regplot(data=combine_complete,
            x='weight',
            y='pick',
            scatter=False,
            ci=None,
            color="navy",
            ax=axs[0,1])
sns.set_style("whitegrid")

sns.scatterplot(data=combine_complete,
                x='bench',
                y='pick',
                hue='pos',
                s=3,
                ax=axs[1,0],
                legend=False)

```

```

axs[1, 0].set(xlabel='bench press reps', ylabel='Pick', title='Bench vs Pick')
sns.regplot(data=combine_complete,
            x='bench',
            y='pick',
            scatter=False,
            ci=None,
            color="navy",
            ax=axs[1,0])
sns.set_style("whitegrid")

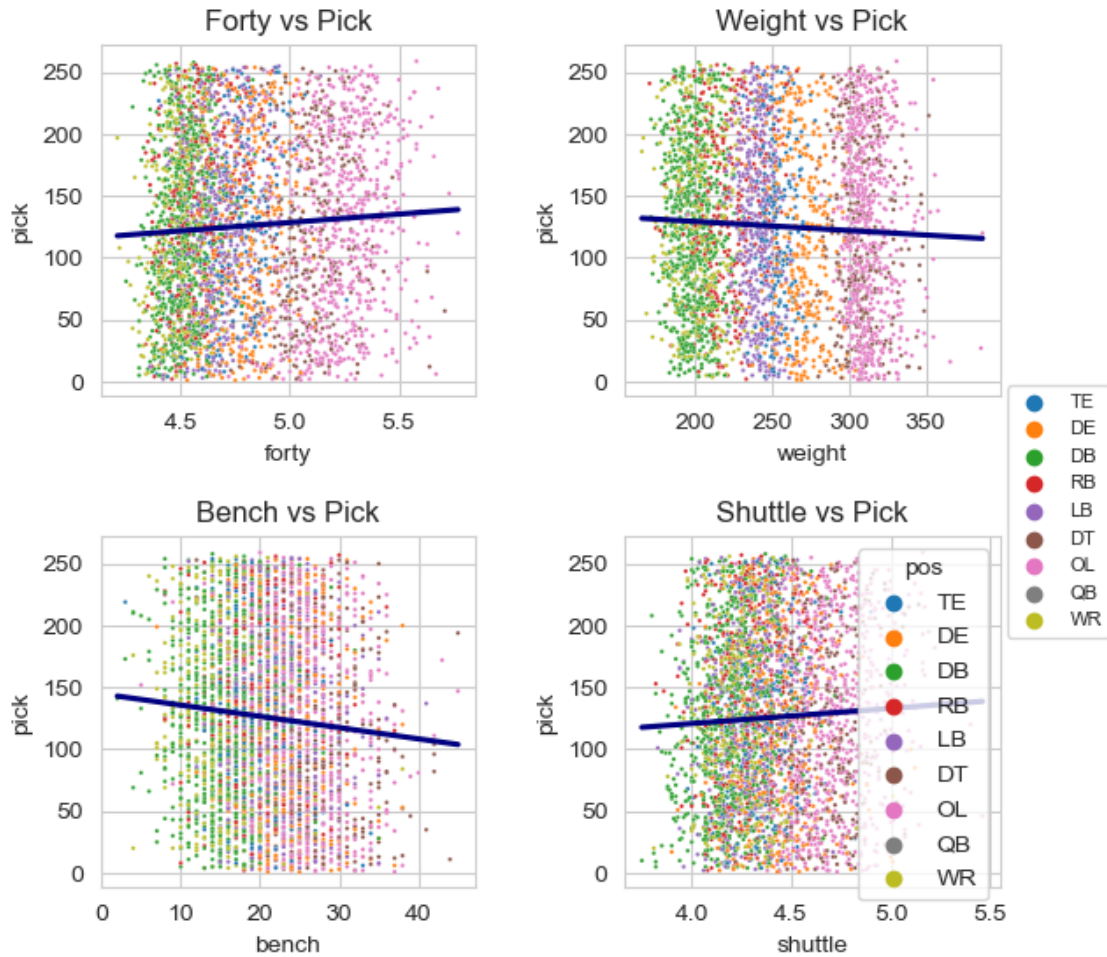
sns.scatterplot(data=combine_complete,
               x='shuttle',
               y='pick',
               hue='pos',
               s=3,
               ax=axs[1,1],
               legend=True)
axs[1, 1].set(xlabel='shuttle time', ylabel='Pick', title='Shuttle vs Pick')

sns.regplot(data=combine_complete,
            x='shuttle',
            y='pick',
            scatter=False,
            ci=None,
            color="navy",
            ax=axs[1,1])
sns.set_style("whitegrid")

handles, labels = axs[1,1].get_legend_handles_labels()
fig.legend(handles=handles, labels=labels, loc='center right', fontsize = 8)
plt.gca().get_legend_handles_labels()
plt.subplots_adjust(left=0.1,
                    bottom=0.5,
                    right=0.9,
                    top=1.5,
                    wspace=0.4,
                    hspace=0.4)

fig.show()

```

[46]: *# This is weight v forty time binned by pick intervals*

```
combine_round_est = combine_complete
combine_round_est['Estimated Round'] = pd.cut(
    combine_round_est['pick'],
    bins=[0, 32, 100, 175, 350],
    labels=['First round', 'Second to Third',
            'Fourth to Fifth', 'Sixth or Later'])
```

```
ax = sns.scatterplot(data=combine_round_est,
                    x='forty',
                    y='weight',
                    hue='Estimated Round',
                    s=10)
```


```
sns.regplot(data=combine_round_est,
            x='forty',
```

```

y='weight',
scatter=False,
ci=None,
color="navy",
ax=ax)

ax.set(xlabel='Forty Yard Dash Time', ylabel='Weight', title='Forty vs Weight')
handles, labels = ax.get_legend_handles_labels()
ax.legend(handles, labels, loc='lower right', fontsize = 8)
sns.set_style("whitegrid")
plt.show()

```



```

[47]: ##### Linear Model w/ model selection #####
# set up basic linear model with complete cases and to simplify further, filter
↳ out undrafted players
model = sm.OLS(y, sm.add_constant(X))
results = model.fit()
print(results.summary())

```

OLS Regression Results

```

=====
Dep. Variable:          pick    R-squared:                0.088
Model:                  OLS     Adj. R-squared:           0.084
Method:                 Least Squares   F-statistic:             18.31
Date:                  Thu, 02 Mar 2023   Prob (F-statistic):      3.50e-56
Time:                  23:27:12   Log-Likelihood:         -19203.
No. Observations:      3418    AIC:                    3.844e+04
Df Residuals:          3399    BIC:                    3.856e+04
Df Model:               18
Covariance Type:       nonrobust
=====

```

```

=====
              coef      std err          t      P>|t|      [0.025      0.975]
-----
const        -233.3898      74.944      -3.114      0.002     -380.330     -86.450
height         1.8195       0.908       2.003      0.045       0.038      3.600
weight        -0.7369       0.117      -6.279      0.000      -0.967     -0.507
hand_size     -3.9821       2.331      -1.708      0.088      -8.553      0.589
arm_length    -5.9567       1.235      -4.824      0.000      -8.378     -3.535
forty         115.4925     10.760     10.733      0.000      94.396     136.589
bench         -0.3985       0.258      -1.543      0.123      -0.905      0.108
vert          -0.2399       0.474      -0.507      0.612      -1.168      0.688
broad_jump    -0.2771       0.210      -1.320      0.187      -0.688      0.134
shuttle       35.6163       8.272       4.306      0.000      19.398     51.834
3cone         0.7358       5.257       0.140      0.889      -9.572     11.044
DB            -30.0442       8.843      -3.398      0.001     -47.382     -12.706
DE            -19.8457     10.113      -1.962      0.050     -39.674      -0.018
DT            -33.2663     11.393      -2.920      0.004     -55.604     -10.928
LB            -16.7721       8.838      -1.898      0.058     -34.101      0.557
OL            -45.8938     12.312      -3.728      0.000     -70.034     -21.754
QB            -45.7701     24.646      -1.857      0.063     -94.093      2.552
RB            -7.6332       8.077      -0.945      0.345     -23.469      8.203
TE            -15.4542     10.342      -1.494      0.135     -35.731      4.822
WR            -18.7102       9.329      -2.006      0.045     -37.002     -0.419
=====

```

```

=====
Omnibus:              448.685    Durbin-Watson:           1.922
Prob(Omnibus):        0.000    Jarque-Bera (JB):       125.679
Skew:                 0.154    Prob(JB):               5.12e-28
Kurtosis:             2.113    Cond. No.:              1.25e+18
=====

```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

[2] The smallest eigenvalue is 1.87e-28. This might indicate that there are strong multicollinearity problems or that the design matrix is singular.

```
[95]: # Backward Elimination to select best model
y = combine_complete['pick']
X = combine_complete.iloc[:, 4:-1]
X = pd.concat([X, pos_one_hot], axis=1)
cols = X.columns

for var in cols:
    model = sm.OLS(y, sm.add_constant(X))
    results = model.fit()
    p_values = pd.Series(results.pvalues, name='pvalues')
    max_p = p_values[p_values==p_values.max()]
    if max_p[0] > 0.05:
        drop_col = max_p.index[0]
        X = X.drop(columns=drop_col)
    else:
        print(results.summary())
        break
```

OLS Regression Results

```
=====
Dep. Variable:          pick    R-squared:                0.084
Model:                  OLS     Adj. R-squared:           0.082
Method:                 Least Squares    F-statistic:        39.27
Date:                  Fri, 03 Mar 2023    Prob (F-statistic):   2.81e-60
Time:                  07:25:40    Log-Likelihood:      -19210.
No. Observations:      3418    AIC:                 3.844e+04
Df Residuals:          3409    BIC:                 3.849e+04
Df Model:               8
Covariance Type:        nonrobust
=====
```

	coef	std err	t	P> t	[0.025	0.975]
const	-222.0558	49.765	-4.462	0.000	-319.628	-124.484
weight	-0.6549	0.088	-7.462	0.000	-0.827	-0.483
arm_length	-6.3900	1.051	-6.079	0.000	-8.451	-4.329
forty	119.8296	9.628	12.445	0.000	100.952	138.708
bench	-0.5226	0.249	-2.102	0.036	-1.010	-0.035
shuttle	37.9831	7.230	5.253	0.000	23.807	52.159
DB	-15.0889	3.667	-4.115	0.000	-22.279	-7.899
DT	-18.5718	5.954	-3.119	0.002	-30.245	-6.898
OL	-28.7441	5.753	-4.996	0.000	-40.024	-17.465

```
=====
Omnibus:                 486.318    Durbin-Watson:           1.925
Prob(Omnibus):            0.000    Jarque-Bera (JB):        129.009
Skew:                     0.147    Prob(JB):                9.68e-29
Kurtosis:                 2.095    Cond. No.:               1.13e+04
=====
```

Notes:

[1] Standard Errors assume that the covariance matrix of the errors is correctly specified.

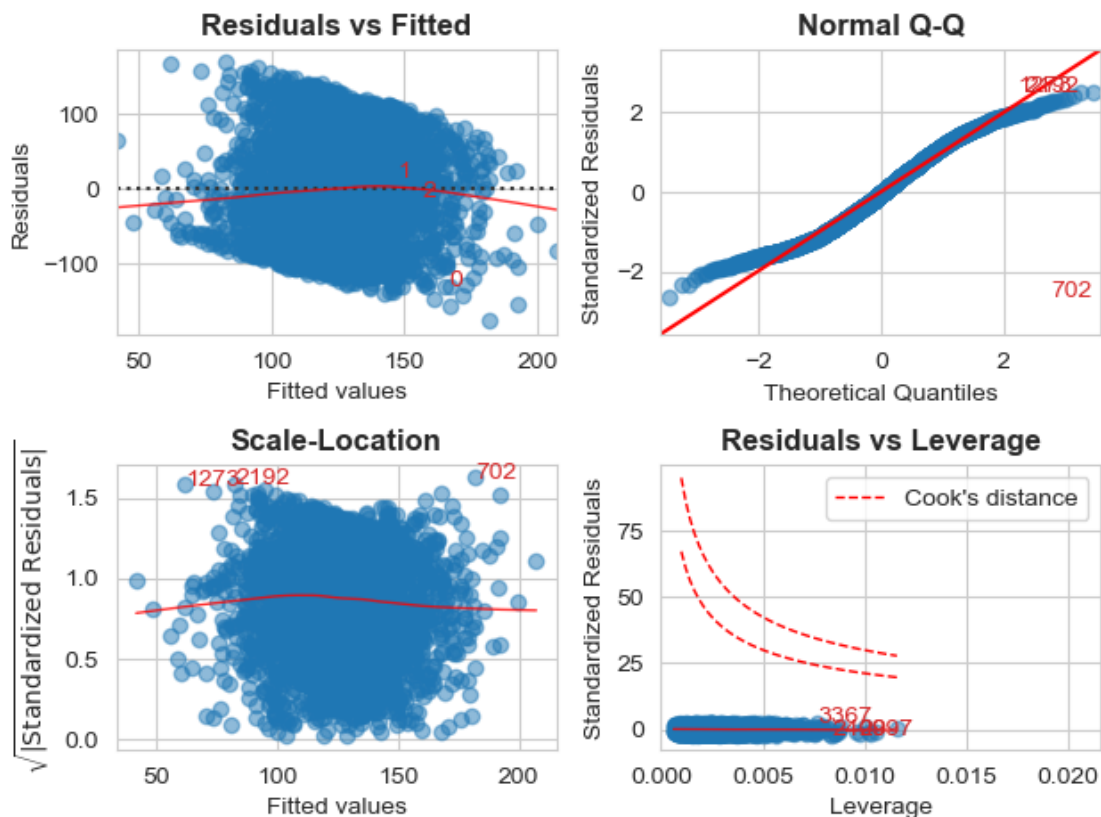
[2] The condition number is large, $1.13\text{e}+04$. This might indicate that there are strong multicollinearity or other numerical problems.

```
[143]: ols = SM.Linear_Reg_Diagnostic(results)
```

```
[156]: fig, axs = plt.subplots(nrows=2, ncols=2)
```

```
ols.residual_plot(ax=axs[0,0])
ols.qq_plot(ax=axs[0,1])
ols.scale_location_plot(ax=axs[1,0])
ols.leverage_plot(ax=axs[1,1])
```

```
fig.tight_layout(pad=1.0)
plt.show()
```



```
[ ]: ### Try robust ###
```

```
[96]: ##### Linear Model #####
# will use the model generated from backward selection above
y = combine_complete['pick']
X = combine_complete.iloc[:, 4:-1]
X = pd.concat([X, pos_one_hot], axis=1)
X = X[['weight', 'arm_length', 'forty', 'bench', 'shuttle', 'DB', 'DT', 'OL']]

cv = KFold(n_splits=3, shuffle=True, random_state=100)
reg = LinearRegression()
scores = cross_val_score(reg, X, y, scoring='neg_mean_squared_error', cv=cv)
rmse = (scores.mean()*-1)**0.5


print(rmse)
```

67.01796200690795

```
[8]: ##### Predictive PCA/PCR #####
y = combine_complete['pick']
X = combine_complete.iloc[:, 4:-1] # do not include position for PCA
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
    random_state=100)


scaler = StandardScaler()
X_train_norm = scaler.fit_transform(X_train.values)
pca = PCA()
pca.fit(X_train_norm)
#transformed = model.transform(X)
features = range(pca.n_components_)

ax = sns.barplot(x=pd.Series(features)+1, y=pd.Series(pca.explained_variance_) /
    pca.explained_variance_.sum())
ax.set(xlabel='PCA Component', ylabel='Variance Proportion', title='PCA Variance_
    Analysis')
sns.set_style("whitegrid")
plt.show()
```



```
[130]: # plot pca vs share of variance:
var_prop = pcr.explained_variance_ / pcr.explained_variance_.sum()
ax = sns.lineplot(x=pd.Series(features), y=pd.Series(var_prop))
ax.set(xlabel='PCA Component', ylabel='Variance Proportion', title='PCA Variance_
↳Analysis')
sns.set_style("whitegrid")
plt.show()

# data shows we should use between 2 and 4 PCs so will use 3 here
```



```
[131]: pca = PCA(n_components=3)
X_train_norm = scaler.fit_transform(X_train.values)
pca.fit_transform(X_train_norm)
linear = pd.DataFrame(pca.components_.T, columns=['PC1', 'PC2', 'PC3'], index=X.
    ↪columns)
linear

#PC1 is generally a size metric, PC2 relates to small hands and short arms, PC3 ↪
    ↪it almost purely strength
```

```
[131]:
```

	PC1	PC2	PC3
height	0.304736	-0.405880	-0.107401
weight	0.385545	-0.038562	0.142756
hand_size	0.228303	-0.461385	0.002039
arm_length	0.263038	-0.533021	-0.253794
forty	0.373551	0.149111	-0.022600
bench	0.232707	-0.048905	0.913638
vert	-0.319665	-0.333223	0.195778
broad_jump	-0.324417	-0.386599	0.068543
shuttle	0.342782	0.121846	-0.135350

3cone 0.344318 0.191882 -0.084137

```
[133]: # Cross-Validated PC Regression
# scale the test set
pca = PCA(n_components=3)
X_norm = scaler.fit_transform(X.values)

# dim reduce X
X_pca = pca.fit_transform(X_norm)
add_pos = pos_one_hot[['DB', 'DT', 'OL']].values
X_pca2 = np.concatenate([X_pca, add_pos], axis=1)
cv = KFold(n_splits=3, shuffle=True, random_state=100)
reg = LinearRegression()

scores = cross_val_score(reg, X_pca2, y, scoring='neg_mean_squared_error', cv=cv)
rmse = (scores.mean()*-1)**0.5


print(rmse)
```

68.05924969610729

```
[134]: # PC scatterplot vs reponse
df_pca = np.concatenate([X_pca, y.values.reshape(-1,1)], axis=1)
df_pca = pd.DataFrame(df_pca, columns=['PC1', 'PC2', 'PC3', 'pick'])

fig, axs = plt.subplots(nrows=3, sharex=True)
sns.regplot(x='pick', y='PC1', data=df_pca, scatter=False, ax=axs[0],
            color="steelblue")
axs[0].set_xlabel('')
sns.regplot(x='pick', y='PC2', data=df_pca, scatter=False, ax=axs[1],
            color="green")
axs[1].set_xlabel('')
sns.regplot(x='pick', y='PC3', data=df_pca, scatter=False, ax=axs[2],
            color="orange")
sns.set_style("whitegrid")
fig.show()

# note how pc1 has little relationship w y and therefore poor mse
```



```
[86]: ##### XG Boost #####
y = combine_complete['pick']
X = combine_complete.iloc[:, 4:-1]
X = pd.concat([X, pos_one_hot], axis=1)
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.3,
    random_state=100)

scaler = StandardScaler()
X_norm = scaler.fit_transform(X)
X_train_norm = scaler.fit_transform(X_train)
X_test_norm = scaler.fit_transform(X_test)

n = 1000
max_depth = 1
xgb_reg = xgb.XGBRegressor(objective='reg:squarederror',
                           n_estimators=n,
                           max_depth=max_depth,
                           eta=0.01,
                           seed=100)
```

```
xgb_reg.fit(X_train_norm, y_train)
preds = xgb_reg.predict(X_test_norm)

rmse = mean_squared_error(preds, y_test, squared=False)
print(rmse)
```

69.51252126008916

```
[8]: # Tune hyperparams w/ gridsearch
param_grid = {
    'max_depth': [1, 2, 3],
    'eta': [0.001, 0.01, 0.1, 0.25, 0.5],
    'n_estimators': [100, 500, 1000],
}
grid_mse = GridSearchCV(estimator=xgb_reg,
                        param_grid=param_grid,
                        scoring='neg_mean_squared_error',
                        cv=3,
                        verbose=1)
mse_grid = grid_mse.fit(X_norm, y)
```

Fitting 3 folds for each of 45 candidates, totalling 135 fits

```
[28]: best_params = grid_mse.best_params_
rmse = np.abs(grid_mse.best_score_)*0.5

print(best_params)
print(rmse)
```

```
{'eta': 0.25, 'max_depth': 1, 'n_estimators': 100}
68.45412247241667
```

```
[29]: xgb_final = xgb.XGBRegressor(objective='reg:squarederror',
                                   n_estimators=100,
                                   max_depth=1,
                                   eta=0.25,
                                   seed=100)

xgb_final.fit(X_norm, y)
```

```
[29]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                   colsample_bylevel=None, colsample_bynode=None,
                   colsample_bytree=None, early_stopping_rounds=None,
                   enable_categorical=False, eta=0.25, eval_metric=None,
                   feature_types=None, gamma=None, gpu_id=None, grow_policy=None,
                   importance_type=None, interaction_constraints=None,
                   learning_rate=None, max_bin=None, max_cat_threshold=None,
```

```

max_cat_to_onehot=None, max_delta_step=None, max_depth=1,
max_leaves=None, min_child_weight=None, missing=nan,
monotone_constraints=None, n_estimators=100, n_jobs=None,
num_parallel_tree=None, predictor=None, ...)

```

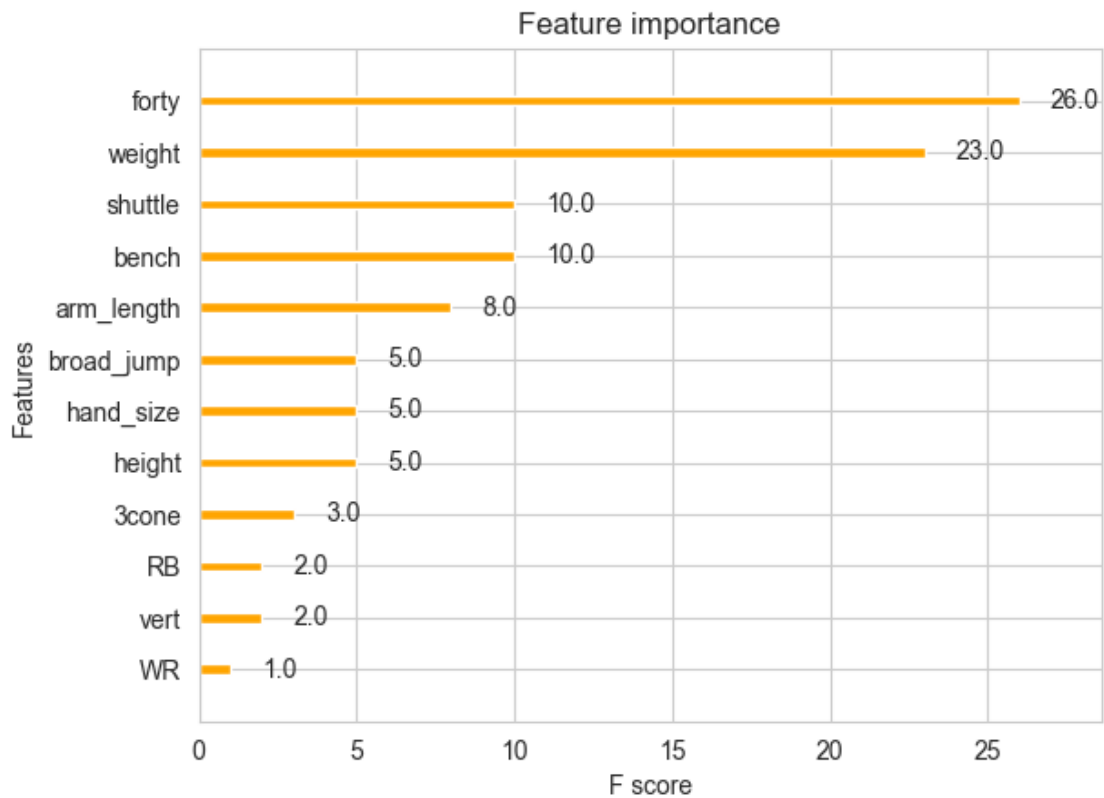
```

[90]: # plot feature importance
feature_names = X.columns

ax = xgb.plot_importance(xgb_final, color="orange")
ylabs = list(ax.get_yticklabels())
dict_features = dict(enumerate(feature_names))
ylabs_stripped = [ylabs[i].get_text().lstrip('f') for i in range(len(ylabs))]
ylabs_stripped = [dict_features[int(i)] for i in ylabs_stripped]

ax.set_yticklabels(ylabs_stripped)
plt.show()

```



```

[ ]: ##### Final Cross-Validated RMSE #####
# Linear: 67.14958397888446
# PCR: 68.68564523905418
# XGBoost: 68.45412247241667

```